01 / THE PROBLEM AND PRODUCT

# KaamProof

Offline AI job closeout for small AC service teams

**Catch missing documentation while the technician is still at the job.**

## The problem

A service visit can leave behind photos, spoken updates and incomplete paperwork. The proposed problem to validate: owners and customers struggle to review what was requested, what the technician reported doing and which required records are missing. A summary alone can make an incomplete record sound complete.

## Who we will build for

Small AC-service businesses and their field technicians. The first prototype covers one predefined service workflow, rather than a full field-service management platform. Customer interviews and demand validation are planned; they have not been completed.

## The product

Technicians open an agreed job checklist, capture before/after photos, and speak short observations. Local AI drafts structured job notes and proposes source links. A checklist engine flags missing required inputs. The technician corrects and approves the report before exporting it.

## A concrete example

A technician says, "Service complete," but has attached only a before photo and no required post-service reading. KaamProof should identify those missing records, request the specific inputs and keep the report pending review. It must not invent a reading or infer performance from a photo.

## What success would look like

A clearer job record with less manual report writing and fewer missing required inputs. We will measure these outcomes; no performance, cost-saving or customer adoption results are claimed at Phase 1.

02 / PHONE-FIRST EXECUTION

# Capture. Check. Review.

A small offline workflow, with people responsible for the final record

## Four screens, one workflow

1. **Job:** scope and required documentation.
2. **Capture:** photos, short audio clips and technician-entered readings.
3. **Review:** missing inputs, suggested questions and source-linked notes.
4. **Report:** human-approved service record and local PDF export.

## Where AI earns its place

The model interprets unstructured observations into proposed job fields and suggests focused clarification questions. It may suggest contradictions for human review. Ordinary application logic enforces required fields and review status. Technician claims, attached evidence and missing information remain distinct.

## Proposed technical approach

React Native interface with a Kotlin module for LiteRT-LM; Gemma 4 E2B is the initial local model candidate. Use local storage for jobs and attachments. Download the model before offline operation. Begin with CPU/GPU inference and verify device compatibility before selecting an acceleration backend.

## Phone hardware and Office Kit

The camera and microphone provide the actual job inputs; local inference makes checking possible without a network. Office Kit will support phone-led development, testing and review using its supported phone/laptop features. The core offline workflow will not depend on laptop inference.

## Deliberate limits

One AC template, short audio clips and a small photo set. English is the first evaluation baseline; Telugu/Hindi are targets pending testing. Typed corrections remain available. No continuous video, automated diagnosis, payment collection, live CRM integration or automatic sending.

**Trust boundary:** documentation completeness is not proof of workmanship, safety or truthful completion. Customer acknowledgement is recorded as acknowledgement only.

# Prove the workflow.

Proposed demonstration, evaluation and component disclosure

## Differentiation we can defend

Voice/photo report generation already exists. Our proposed focus is on-device completeness checking before the technician leaves, with statements linked to supplied sources. CompanyCam documents deferred processing for offline walkthroughs; MaintainX offers AI voice summaries. This is a focused workflow distinction, not a claim of global uniqueness.

## Planned live demonstration

Provision the model, then enable airplane mode. Open a service job, capture an incomplete record and show missing required evidence. Supply the missing inputs, review the draft and export the report on the phone. The product is currently a concept; this demonstration has not yet been implemented.

## Evaluation and build priorities

First verify one local extraction on the target Android device. Then build capture, persistence, checklist checks and export. Evaluate complete, incomplete and contradictory records. Measure field extraction accuracy, gap-detection precision/recall, unsupported claims, report time versus manual entry, latency and recovery after restart. All results remain to be measured.

## Relevant experience and main risks

The team leader's recorded prior work includes React/Python applications, voice and cloud-LLM integrations, and a document-processing prototype. These are background experience, not KaamProof capabilities. Main risks: Android integration, noisy speech, unreliable source matching and technician adoption. A real service owner must validate the checklist before broader rollout.

## Originality and pre-existing components

This document describes our proposed KaamProof workflow. Planned external components include React Native, Kotlin, LiteRT-LM and pretrained Gemma 4 weights. No custom model training or completed KaamProof application is claimed. Any reused templates, libraries or prior project code will be disclosed, with applicable licences and event rules followed. Global uniqueness has not been established.

**Selected references**

Organiser: iqoo.reskilll.com/guide
Runtime: Google LiteRT-LM Android documentation
Model: Google Gemma 4 mobile deployment
Comparison: CompanyCam AI documents; MaintainX AI